

**Roll No:** 20BCE204

**Name:** Dhyan Patel

**Course:** Cloud Computing

**Practical No:** 5

**Aim:** Working with EBS

***EBS – Elastic block store***, a key underlying storage mechanism for Amazon EC2 instances.

Amazon Elastic Block Store (EBS) is a scalable block storage service offered by Amazon Web Services (AWS). It is designed to provide persistent block-level storage volumes for use with Amazon Elastic Compute Cloud (EC2) instances. EBS volumes are highly available, reliable, and durable storage systems that can be attached to EC2 instances as virtual hard drives.

EBS volumes are stored in Amazon's redundant storage infrastructure, which provides high durability and availability. EBS volumes can be created and deleted independently of EC2 instances, and can be moved between instances. They are also designed to be resilient to failures, and can automatically recover from hardware or software failures.

EBS is a reliable and scalable storage solution for EC2 instances, providing persistent block-level storage that is highly available and durable.

EBS volumes are available in two types: solid-state drives (SSDs) and hard disk drives (HDDs).

Now, let's start the lab

Here, we are already given with an instance name Lab. Now, let's create a volume from the EC2 dashboard.

Here to create our volume we need to specify our volume type, size of volume and availability zone. After that we will be assign a tag by providing a key value pair to this volume.

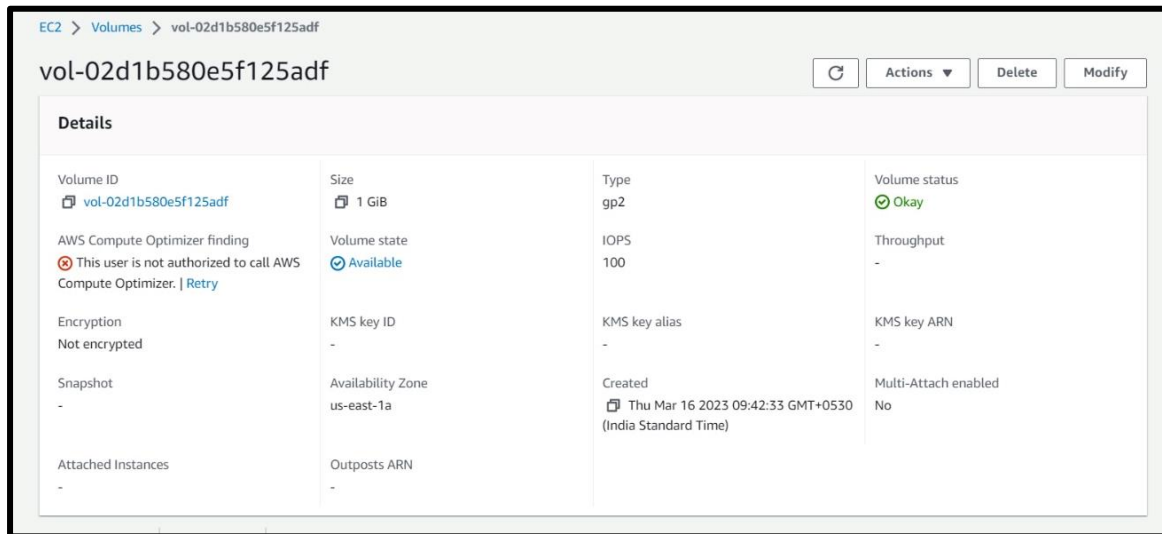
Now, click on **Create Volume** button.

After creating the volume, it will take few minutes to move from **Creating state** to **available state**.

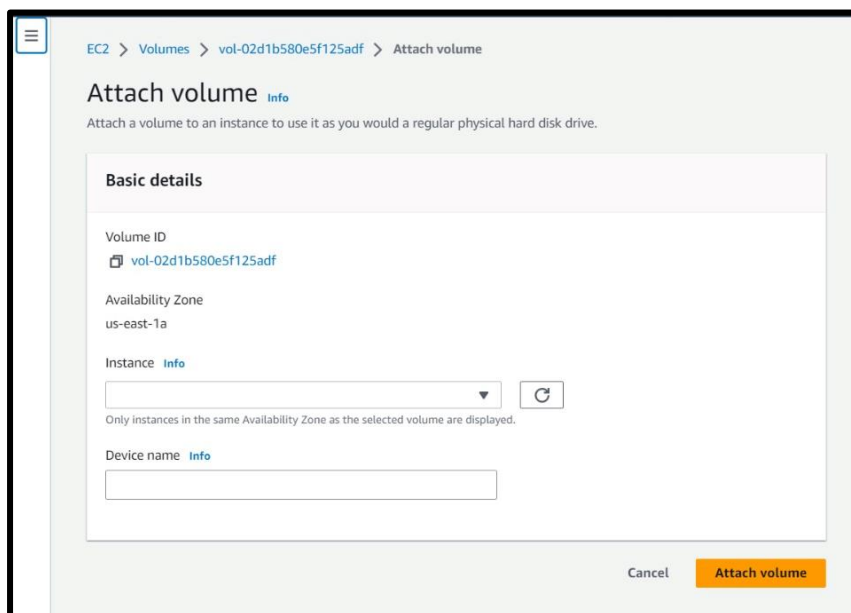
Throughput	Snapshot	Created	Availability Zone	Volume state	Alarm status	Attached Instances	Volume
-	snap-0e1f706...	2023/03/16 09:39 GMT+5:...	us-east-1a	In-use	No alarms	+ i-022774e6e6026482d (La...	Okay
-	snap-0e1f706...	2023/03/16 09:41 GMT+5:...	us-east-1a	In-use	No alarms	+ i-0dc6df8bdf3b220cd (Bas...	Okay
-	-	2023/03/16 09:42 GMT+5:...	us-east-1a	Available	No alarms	+ -	Okay

Here, our volume is now available after refreshing the page.

Let's now attach this volume to our EC2 instance name LAB.  
For that, open the volume that we have just created.



Click on **Action** and then on **Attach Volume**.




Now choose the instance for which we need to use this volume.  
Also specify the device name as **`/dev/sdf`** and then click on the **Attach Volume** button


### Attach volume [Info](#)

Attach a volume to an instance to use it as you would a regular physical hard disk drive.

#### Basic details

Volume ID  
 [vol-02d1b580e5f125adf](#)


Availability Zone  
us-east-1a

Instance [Info](#)  
 

Only instances in the same Availability Zone as the selected volume are displayed.







Device name [Info](#)

Recommended device names for Linux: /dev/sda1 for root volume, /dev/sd[f-p] for data volumes.



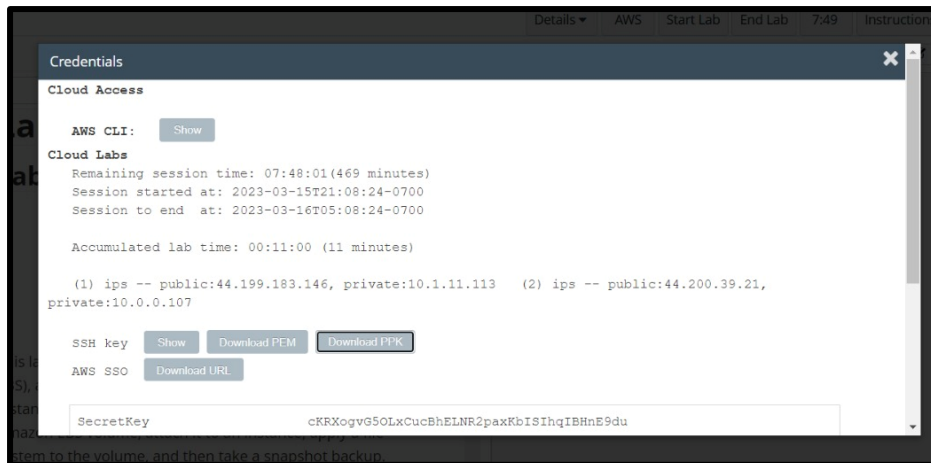
Newer Linux kernels may rename your devices to `/dev/xvdf` through `/dev/xvdp` internally, even when the device name entered here (and shown in the details) is `/dev/sdf` through `/dev/sdp`.

Cancel
Attach volume

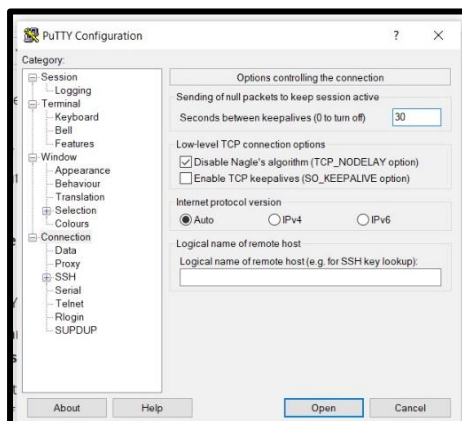
vol-02d1b580e5f125adf			
Details			
Volume ID  <a href="#">vol-02d1b580e5f125adf</a>	Size  1 GiB	Type gp2	Volume status  <span>Okay</span>
AWS Compute Optimizer finding  This user is not authorized to call AWS Compute Optimizer. <a href="#">Retry</a>	Volume state  <span>In-use</span>	IOPS 100	Throughput -
Encryption Not encrypted	KMS key ID -	KMS key alias -	KMS key ARN -
Snapshot -	Availability Zone us-east-1a	Created  Thu Mar 16 2023 09:42:33 GMT+0530 (India Standard Time)	Multi-Attach enabled No
Attached Instances <a href="#">i-022774e6e6026482d (Lab)</a> : /dev/sdf (attached)	Outposts ARN -		

Now, we will mount this volume on over EC2 instance named LAB from our PC using **Putty**.

For that we will require the credential, to get that we will click on **details** button which is available on the Lab manual page. From their we will download the ppk credentials.



Now open putty, go to **Connections** and in that change the **Seconds between keepalives** to **30**.



Now copy the **public DNS or IPv4 address** of our instance name LAB.

Configure the putty using this public DNS and also authenticate with the key which we have downloaded and start the session.



```
ec2-user@ip-10-1-11-113:~  
login as: ec2-user  
Authenticating with public key "imported-openssh-key"  
  
 _ | _ | _ )  
 _ | ( _ | /  Amazon Linux 2 AMI  
 _ | \ _ | _ |  
  
https://aws.amazon.com/amazon-linux-2/  
[ec2-user@ip-10-1-11-113 ~]$
```

Now let's view the storage available on your instance by write the command **df -h**

```
ec2-user@ip-10-1-11-113:~  
login as: ec2-user  
Authenticating with public key "imported-openssh-key"  
  
 _ | _ | _ )  
 _ | ( _ | /  Amazon Linux 2 AMI  
 _ | \ _ | _ |  
  
https://aws.amazon.com/amazon-linux-2/  
[ec2-user@ip-10-1-11-113 ~]$ df -h  
Filesystem      Size  Used Avail Use% Mounted on  
devtmpfs        484M   0  484M   0% /dev  
tmpfs           492M   0  492M   0% /dev/shm  
tmpfs           492M 412K  491M   1% /run  
tmpfs           492M   0  492M   0% /sys/fs/cgroup  
/dev/xvda1      8.0G  1.5G   6.5G  19% /  
tmpfs           99M   0   99M   0% /run/user/1000  
[ec2-user@ip-10-1-11-113 ~]$
```

Our new volume is not yet shown.

First of all, we need to create a file system. For example, let's create a file system named **ext3**

```
ec2-user@ip-10-1-11-113:~  
login as: ec2-user  
Authenticating with public key "imported-openssh-key"  
  
 _ _ | _ _ | _ )  
 _ | ( _ _ | /  Amazon Linux 2 AMI  
 _ | \ _ _ | _ _ |  
  
https://aws.amazon.com/amazon-linux-2/  
[ec2-user@ip-10-1-11-113 ~]$ df -h  
Filesystem      Size  Used Avail Use% Mounted on  
devtmpfs        484M   0  484M   0% /dev  
tmpfs           492M   0  492M   0% /dev/shm  
tmpfs           492M  412K  491M   1% /run  
tmpfs           492M   0  492M   0% /sys/fs/cgroup  
/dev/xvda1      8.0G  1.5G   6.5G  19% /  
tmpfs           99M    0   99M   0% /run/user/1000  
[ec2-user@ip-10-1-11-113 ~]$ sudo mkfs -t ext3 /dev/sdf  
mke2fs 1.42.9 (28-Dec-2013)  
Filesystem label=  
OS type: Linux  
Block size=4096 (log=2)  
Fragment size=4096 (log=2)  
Stride=0 blocks, Stripe width=0 blocks  
65536 inodes, 262144 blocks  
13107 blocks (5.00%) reserved for the super user  
First data block=0  
Maximum filesystem blocks=268435456  
8 block groups  
32768 blocks per group, 32768 fragments per group  
8192 inodes per group  
Superblock backups stored on blocks:  
    32768, 98304, 163840, 229376  
  
Allocating group tables: done  
Writing inode tables: done  
Creating journal (8192 blocks): done  
Writing superblocks and filesystem accounting information: done  
  
[ec2-user@ip-10-1-11-113 ~]$
```

Now we will create a directory to mount our volume. Let's say we have created a directory ***/mnt/data-store***.

Now we will mount our ***/dev/sdf*** volume to this directory ***/mnt/data-store***

```
[ec2-user@ip-10-1-11-113 ~]$ sudo mkdir /mnt/data-store  
[ec2-user@ip-10-1-11-113 ~]$ sudo mount /dev/sdf /mnt/data-store  
[ec2-user@ip-10-1-11-113 ~]$ echo "/dev/sdf /mnt/data-store ext3 defaults,noatime 1 2" | sudo tee -a /etc/fstab  
/dev/sdf /mnt/data-store ext3 defaults,noatime 1 2  
[ec2-user@ip-10-1-11-113 ~]$ cat /etc/fstab  
#  
UUID=c6fb2880-7e25-4bc8-b8fe-dc604e20c5da / xfs defaults,noatime 1 1  
/dev/sdf /mnt/data-store ext3 defaults,noatime 1 2  
[ec2-user@ip-10-1-11-113 ~]$
```

Now again let's see the storage of our EC2 instance

```
[ec2-user@ip-10-1-11-113 ~]$ df -h  
Filesystem      Size  Used Avail Use% Mounted on  
devtmpfs        484M   0  484M   0% /dev  
tmpfs           492M   0  492M   0% /dev/shm  
tmpfs           492M  412K  491M   1% /run  
tmpfs           492M   0  492M   0% /sys/fs/cgroup  
/dev/xvda1      8.0G  1.5G   6.5G  19% /  
tmpfs           99M    0   99M   0% /run/user/1000  
/dev/xvdf       975M  60K  924M   1% /mnt/data-store  
[ec2-user@ip-10-1-11-113 ~]$
```

As we can see an additional line is present **/dev/xvdf** , which our volume which we just mount on our EC2 instance.

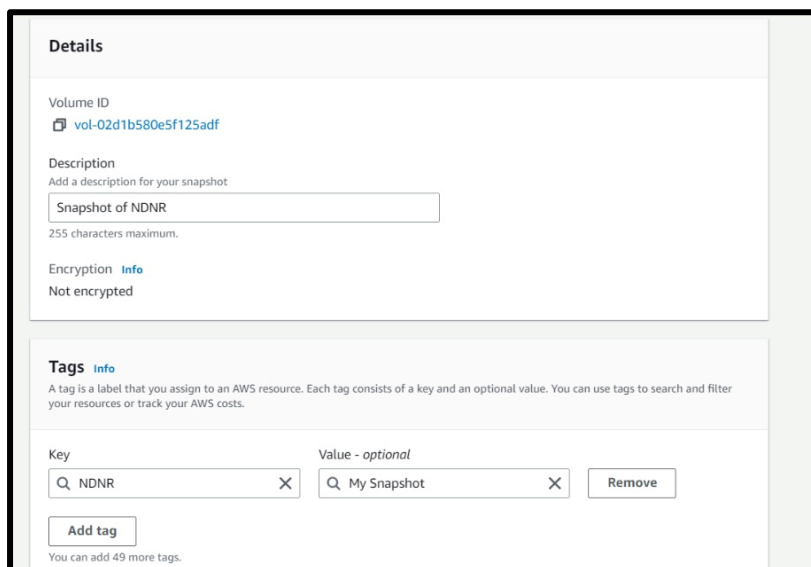
Let's create a file named as **file.txt** and write some content inside that file.

```
[ec2-user@ip-10-1-11-113 ~]$ sudo sh -c "echo This is the group lab of NDNR > /mnt/data-store/file.txt"
[ec2-user@ip-10-1-11-113 ~]$ cat /mnt/data-store/file.txt
This is the group lab of NDNR
[ec2-user@ip-10-1-11-113 ~]$
```

Next, we will create the snapshot of our volume.

To create snapshot, we have to go to our volume and in action we will find option of creating snapshot of the volume, on click we will find below dialogue box.

Here, we can add the description of our snapshot. We also have to add a tag which is key value pair.



The screenshot shows the 'Create Snapshot' dialog box in the AWS console. It is divided into two main sections: 'Details' and 'Tags'.

**Details:**

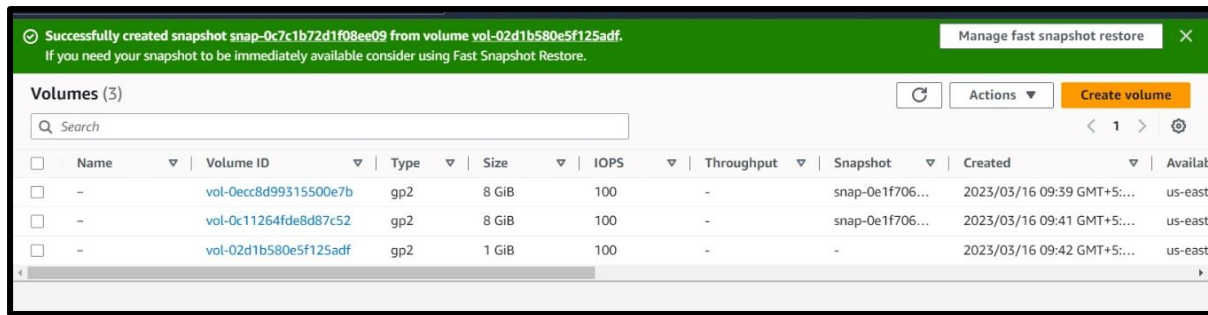
- Volume ID:** vol-02d1b580e5f125adf
- Description:** A text input field containing 'Snapshot of NDNR'. Below it, it says '255 characters maximum.'
- Encryption:** Not encrypted (with an 'Info' link).

**Tags:**

- A description: 'A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to search and filter your resources or track your AWS costs.'
- Key:** A text input field containing 'Q NDNR'.
- Value - optional:** A text input field containing 'Q My Snapshot'.
- Buttons:** 'Add tag', 'Remove', and 'Create snapshot' (partially visible).
- Footer:** 'You can add 49 more tags.'



After creating the snapshot, we will get a popup message of successfully creation of snapshot of our volume.





Now we will go to Snapshots which will be available in the left navigation plane of the EC2 dashboard.

Inside we will find a snapshot will be existing, which the one we created just now. First, it will be in ***pending state*** and after few minutes the creation of snapshot will be completed and it will change to ***completed state***



Now, we will remove our file.txt from our volume using ***rm*** command.

```
[ec2-user@ip-10-1-11-113 ~]$ sudo rm /mnt/data-store/file.txt
[ec2-user@ip-10-1-11-113 ~]$ ls /mnt/data-store/
lost+found
[ec2-user@ip-10-1-11-113 ~]$
```

Now if we want to retrieve our ***file.txt***, is it possible? how?


- ⇒ Yes, it is possible
- ⇒ We can retrieve our file.txt from our snapshot because the snapshot will contain all the files of the volume.
- ⇒ How to do that
  - First, we have to create a volume from the snapshot

- Then we have to attach this volume to our EC2 instance
- After attaching, we will mount the volume to our EC2 instance and then we will be available to access our ***file.txt***

**Create volume** [Info](#)

Create an Amazon EBS volume to attach to any EC2 instance in the same Availability Zone.

**Volume settings**

Snapshot ID  
 snap-0c7c1b72d1f08ee09

Volume type [Info](#)  
 General Purpose SSD (gp2) ▼

Size (GiB)  
  
 Min: 1 GiB, Max: 16384 GiB. The value must be an integer.

IOPS  
 100 / 3000  
 Baseline of 3 IOPS per GiB with a minimum of 100 IOPS, burstable to 3000 IOPS.

Throughput (MiB/s) [Info](#)  
 Not applicable

Availability Zone [Info](#)  
 us-east-1a ▼

Fast snapshot restore [Info](#)  
☒ Not enabled for selected snapshot

At the time of creating volume from snapshot we can change the size of our volume disc and also the type.

We also have to provide tag which a key value pair.

**Tags - optional** [Info](#)

A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to search and filter your resources or track your AWS costs.

Key Value - optional

You can add 49 more tags.

✓ Successfully created volume vol-0e6c992ffc698425.

**Snapshots** [Refresh](#) [Recycle Bin](#) [Actions](#) [Create snapshot](#)

Owned by me ▼

	Name	Snapshot ID	Size	Description	Storage...	Snapshot status	Started
Loading snapshots							

Now we will attach this volume to our EC2 instance named LAB.

**Basic details**

Volume ID  
vol-0e6c992ffc698425

Availability Zone  
us-east-1a

Instance [Info](#)  
i-022774e6e6026482d

Only instances in the same Availability Zone as the selected volume are displayed.

Device name [Info](#)  
/dev/sdg

Recommended device names for Linux: /dev/sda1 for root volume, /dev/sd[f-p] for data volumes.

Newer Linux kernels may rename your devices to /dev/xvdf through /dev/xvdp internally, even when the device name entered here (and shown in the details) is /dev/sdf through /dev/sdp.

Cancel **Attach volume**

Successfully attached volume vol-0e6c992ffc698425 to instance i-022774e6e6026482d.

**Volumes (4)**

	Name	Volume ID	Type	Size	IOPS	Throughput	Snapshot	Created	Availat
<input type="checkbox"/>	-	vol-0ecc8d99315500e7b	gp2	8 GiB	100	-	snap-0e1f706...	2023/03/16 09:39 GMT+5:...	us-east
<input type="checkbox"/>	-	vol-0c11264fde8d87c52	gp2	8 GiB	100	-	snap-0e1f706...	2023/03/16 09:41 GMT+5:...	us-east
<input type="checkbox"/>	-	vol-02d1b580e5f125adf	gp2	1 GiB	100	-	-	2023/03/16 09:42 GMT+5:...	us-east
<input type="checkbox"/>	-	vol-0e6c992ffc698425	gp2	1 GiB	100	-	snap-0c7c1b7...	2023/03/16 10:28 GMT+5:...	us-east

Now we will mount this volume on our EC2 instance in a directory ***/mnt/data-store2***

```
THIS IS THE GROUP LAB OF NDNR
[ec2-user@ip-10-1-11-113 ~]$ sudo rm /mnt/data-store/file.txt
[ec2-user@ip-10-1-11-113 ~]$ ls /mnt/data-store/
lost+found
[ec2-user@ip-10-1-11-113 ~]$ sudo mkdir /mnt/data-store2
[ec2-user@ip-10-1-11-113 ~]$ sudo mount /dev/sdg /mnt/data-store2
[ec2-user@ip-10-1-11-113 ~]$ ls /mnt/data-store2/
file.txt  lost+found
[ec2-user@ip-10-1-11-113 ~]$ cat /mnt/data-store2/file.txt
This is the group lab of NDNR
[ec2-user@ip-10-1-11-113 ~]$
```

As we can see that **file.txt** is not available inside the **data-store** but is available inside the **data-store2**. This is because we have mounted the volume which is the snapshot of the previous volume.