

Roll No: 20BCE204

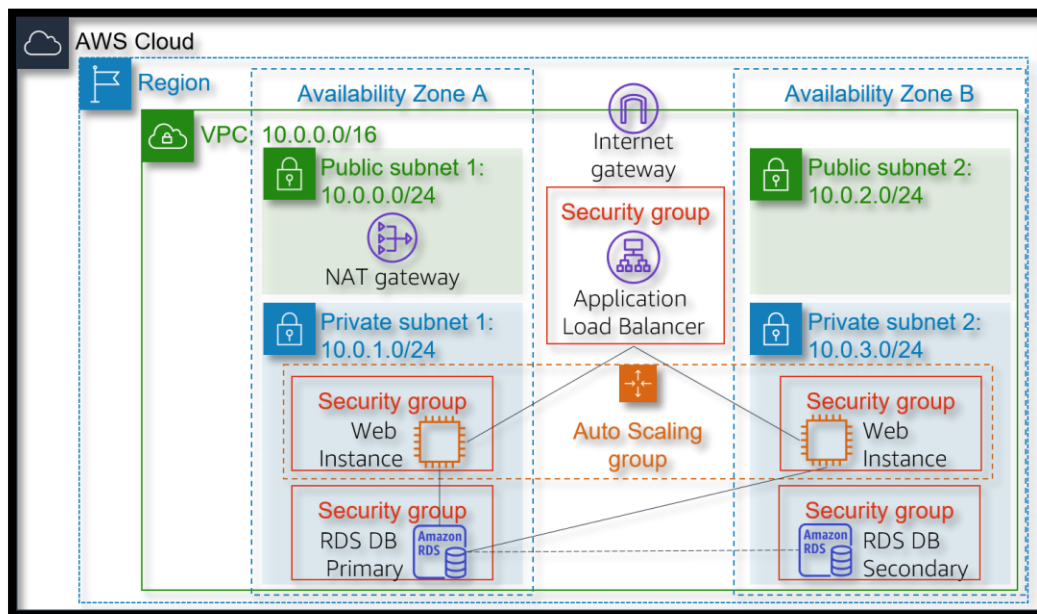
Name: Dhyan Patel

Course: Cloud Computing

Practical No: 7

Aim: Working with an IaaS Cloud Computing: Using AWS (Amazon Web Services) to understating the following concept.Auto-scaling in amazon

Auto Scaling helps you maintain application availability and allows you to scale your Amazon EC2 capacity out or in automatically according to conditions you define. Auto Scaling can also automatically increase the number of Amazon EC2 instances during demand spikes to maintain performance and decrease capacity during lulls to reduce costs. Auto Scaling is well suited to applications that have stable demand patterns or that experience hourly, daily, or weekly variability in usage.



Task 1: Create a Launch Configuration and an Auto Scaling Group

Amazon EC2 Auto Scaling

helps maintain the availability of your applications

Auto Scaling groups are collections of Amazon EC2 instances that enable automatic scaling and fleet management features. These features help you maintain the health and availability of your applications.

Create Auto Scaling group

Get started with EC2 Auto Scaling by creating an Auto Scaling group.

Create Auto Scaling group

1. In the left navigation pane, click Launch Configurations.
2. Click Create launch configuration

Step 1

Choose launch template or configuration

Step 2

Choose instance launch options

Step 3 - optional

Configure advanced options

Step 4 - optional

Configure group size and scaling policies

Step 5 - optional

Add notifications

Step 6 - optional

Add tags

Step 7

Choose launch template or configuration [Info](#)

Specify a launch template that contains settings common to all EC2 instances that are launched by this Auto Scaling group. If you currently use launch configurations, you might consider migrating to launch templates.

Name

Auto Scaling group name

Enter a name to identify the group.

autoscaled

Must be unique to this account in the current Region and no more than 255 characters.

Launch template [Info](#)

[Switch to launch configuration](#)

Launch template

Choose a launch template that contains the instance-level settings, such as the Amazon Machine Image (AMI), instance type, key pair, and security groups.

Select a launch template

Create a launch template [↗](#)

3. Configure these settings:
 - Launch configuration name: LabConfig
 - Amazon Machine Image (AMI) Choose *Web Server AMI*

Create launch template

Creating a launch template allows you to create a saved instance configuration that can be reused, shared and launched at a later time. Templates can have multiple versions.

Launch template name and description

Launch template name - *required*

Must be unique to this account. Max 128 chars. No spaces or special characters like '&', '*', '@'.

Template version description

Max 255 chars

Auto Scaling guidance [Info](#)

Select this if you intend to use this template with EC2 Auto Scaling

☒ Provide guidance to help me set up a template that I can use with EC2 Auto Scaling

- Instance type: Select *t3.micro*
 - Choose Select an existing security group
4. Under Key pair configure:
- Key pair options: *Choose an existing key pair*
 - Existing key pair: vockey

▼ Instance type
Info
Advanced

Instance type

t2.micro
Free tier eligible

Family: t2 1 vCPU 1 GiB Memory
On-Demand Windows pricing: 0.0162 USD per Hour
On-Demand SUSE pricing: 0.0116 USD per Hour
On-Demand RHEL pricing: 0.0716 USD per Hour
On-Demand Linux pricing: 0.0116 USD per Hour

☐ All generations

Compare instance types

▼ Key pair (login)
Info

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name

vockey

Create new key pair

▼ Network settings
Info

Subnet
Info

Don't include in launch template

Create new subnet

When you specify a subnet, a network interface is automatically added to your template.

Firewall (security groups)
Info

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

☒ Select existing security group

☐ Create security group

Security groups
Info

Select security groups

Web Security Group sg-0fc56e4d561470ef2 X
VPC: vpc-0c38fe58f6d15ea6e

Compare security group rules

▶ Advanced network configuration

- Select I acknowledge...
- Click Create launch configuration

You will now create an Auto Scaling group that uses this Launch Configuration.

5. Select the checkbox for the *LabConfig* Launch Configuration.
6. From the Actions menu, choose *Create Auto Scaling group*
7. Enter Auto Scaling group name:
 - Name: Lab Auto Scaling Group
8. Choose Next
9. On the Network page configure

Network: *Lab VPC*

Choose instance launch options [Info](#)

Choose the VPC network environment that your instances are launched into, and customize the instance types and purchase options.

Network [Info](#)

For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.

VPC

Choose the VPC that defines the virtual network for your Auto Scaling group.

vpc-0b53808eb23040f52
172.31.0.0/16 Default

▼

↺

[Create a VPC](#) [↗](#)

▼ Application and OS Images (Amazon Machine Image) - required [Info](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

 Search our full catalog including 1000s of application and OS images


Recents | Quick Start

 Currently in use


[Browse more AMIs](#)
Including AMIs from
AWS, Marketplace and
the Community

Amazon Machine Image (AMI)

amzn2-ami-hvm-2.0.20230404.1-x86_64-gp2
ami-05f408238af346b4f
2023-04-11T07:52:19.000Z architecture: 64-bit (x86) Virtualization: hvm ENA enabled: true
Root device type: ebs




Description

Amazon Linux 2 AMI 2.0.20230404.1 x86_64 HVM gp2

Architecture	AMI ID	
x86_64	ami-05f408238af346b4f	Verified provider

Availability Zones and subnets
Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.

Select Availability Zones and subnets ▼ 

us-east-1a | subnet-0aba284e8cb7239ef ✕
172.31.0.0/20 Default


us-east-1b | subnet-0ee74ed220509b420 ✕
172.31.80.0/20 Default

us-east-1c | subnet-018e7d63786eb459e ✕
172.31.16.0/20 Default

us-east-1d | subnet-0fa37a0c33b896943 ✕
172.31.32.0/20 Default

us-east-1e | subnet-05249d9b43c943c91 ✕
172.31.48.0/20 Default

us-east-1f | subnet-0345eef417871de4d ✕
172.31.64.0/20 Default

[Create a subnet](#) 

10. Choose Next
11. In the Load balancing - *optional* pane, choose Attach to an existing load balancer
12. In the Attach to an existing load balancer pane, use the dropdown list to select *LabGroup*.

Configure advanced options - *optional* [Info](#)

Integrate your Auto Scaling group with other services to distribute network traffic across multiple servers using a load balancer or to establish service-to-service communications using VPC Lattice. You can also set options that give you more control over health check replacements and monitoring.

Load balancing [Info](#)

Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.

☐ No load balancer
Traffic to your Auto Scaling group will not be fronted by a load balancer.

☐ Attach to an existing load balancer
Choose from your existing load balancers.

☒ Attach to a new load balancer
Quickly create a basic load balancer to attach to your Auto Scaling group.

Attach to a new load balancer

Define a new load balancer to create for attachment to this Auto Scaling group.

Load balancer type

Choose from the load balancer types offered below. Type selection cannot be changed after the load balancer is created. If you need a different type of load balancer than those offered here, [visit the Load Balancing console](#). [↗](#)

Attach to a new load balancer

Define a new load balancer to create for attachment to this Auto Scaling group.

Load balancer type

Choose from the load balancer types offered below. Type selection cannot be changed after the load balancer is created. If you need a different type of load balancer than those offered here, [visit the Load Balancing console](#). [↗](#)

☒ Application Load Balancer
HTTP, HTTPS

☐ Network Load Balancer
TCP, UDP, TLS

Load balancer name

Name cannot be changed after the load balancer is created.

autoscale-1

Load balancer scheme

Scheme cannot be changed after the load balancer is created.

☐ Internal

☒ Internet-facing

13. Choose Next

Listeners and routing

If you require secure listeners, or multiple listeners, you can configure them from the [Load Balancing console](#) after your load balancer is created.

Protocol	Port	Default routing (forward to)
HTTP	80	Create a target group ▼

New target group name
An instance target group with default settings will be created.

autoscale-1

Tags - optional
Consider adding tags to your load balancer. Tags enable you to categorize your AWS resources so you can more easily manage them.

Add tag

50 remaining

14. Under Group size, configure:
- Desired capacity: 2
 - Minimum capacity: 2
 - Maximum capacity: 4

Configure group size and scaling policies - optional [Info](#)

Set the desired, minimum, and maximum capacity of your Auto Scaling group. You can optionally add a scaling policy to dynamically scale the number of instances in the group.

Group size - optional [Info](#)

Specify the size of the Auto Scaling group by changing the desired capacity. You can also specify minimum and maximum capacity limits. Your desired capacity must be within the limit range.

Desired capacity

2

Minimum capacity


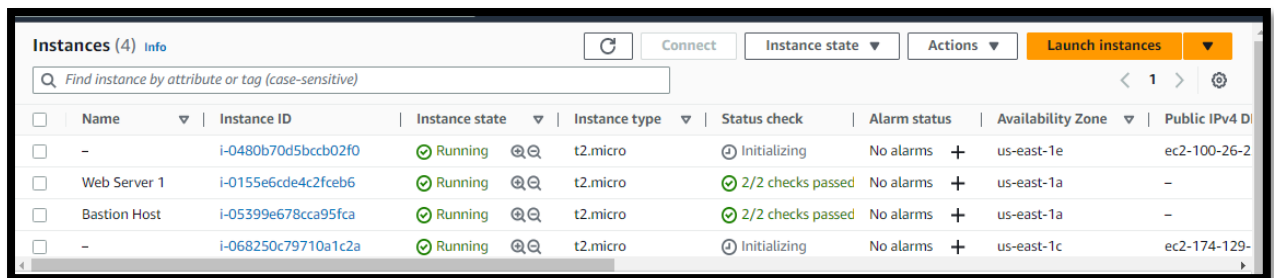
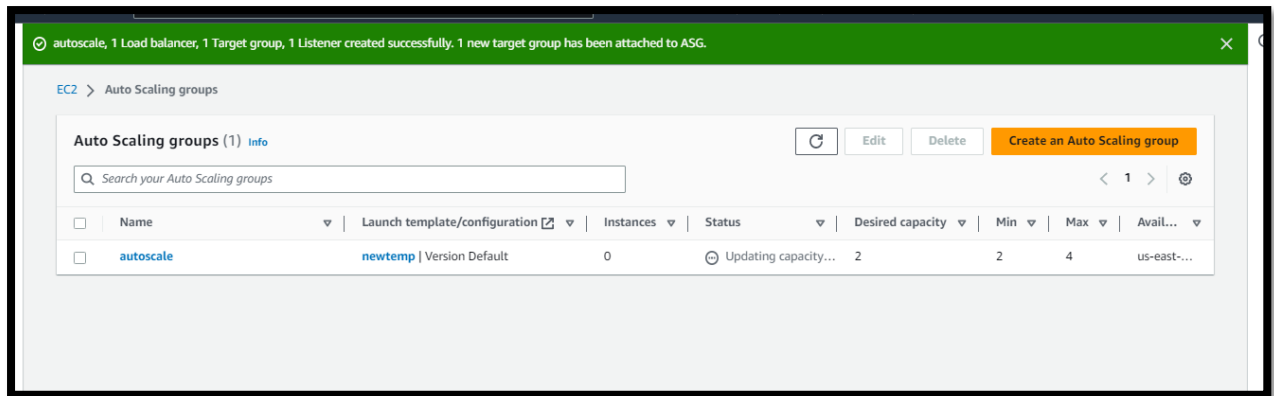
2

Maximum capacity

4

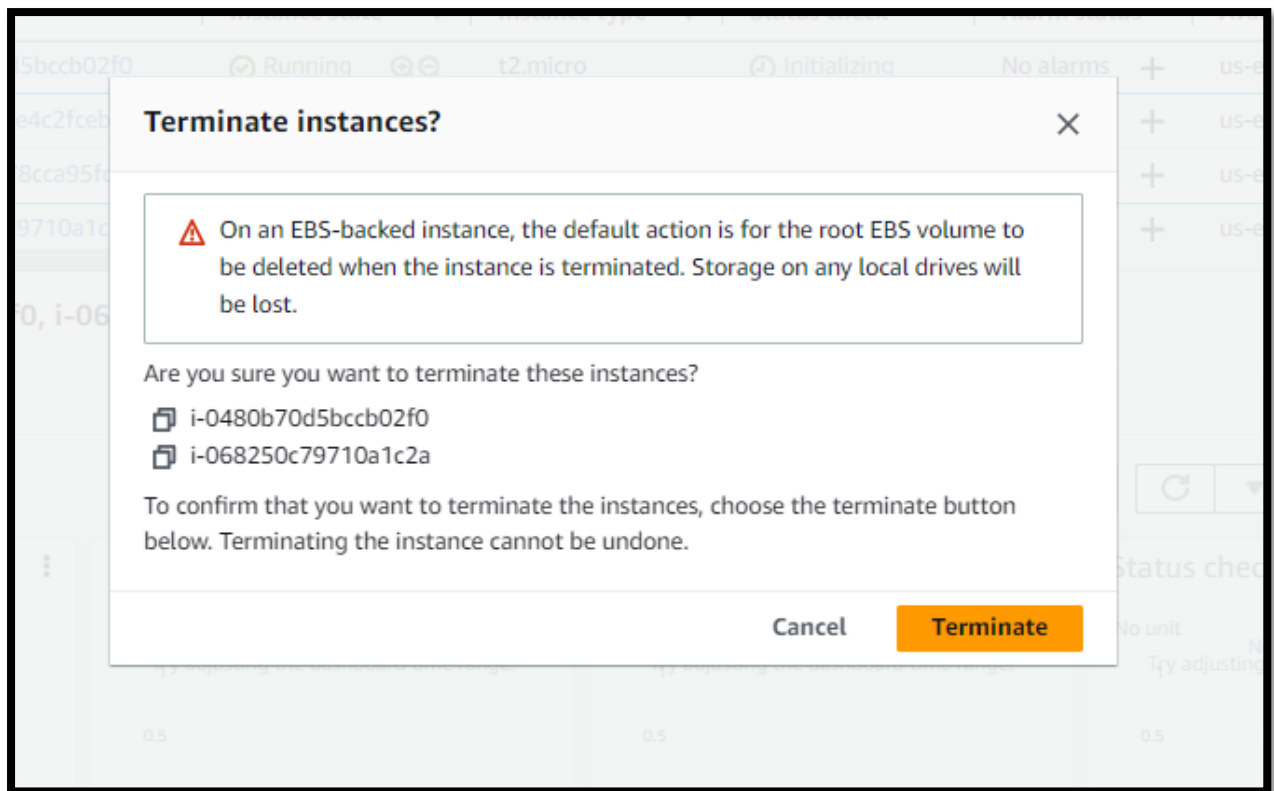
15. Choose Next

Your Auto Scaling group will initially show an instance count of zero, but new instances will be launched to reach the Desired count of 2 instances.



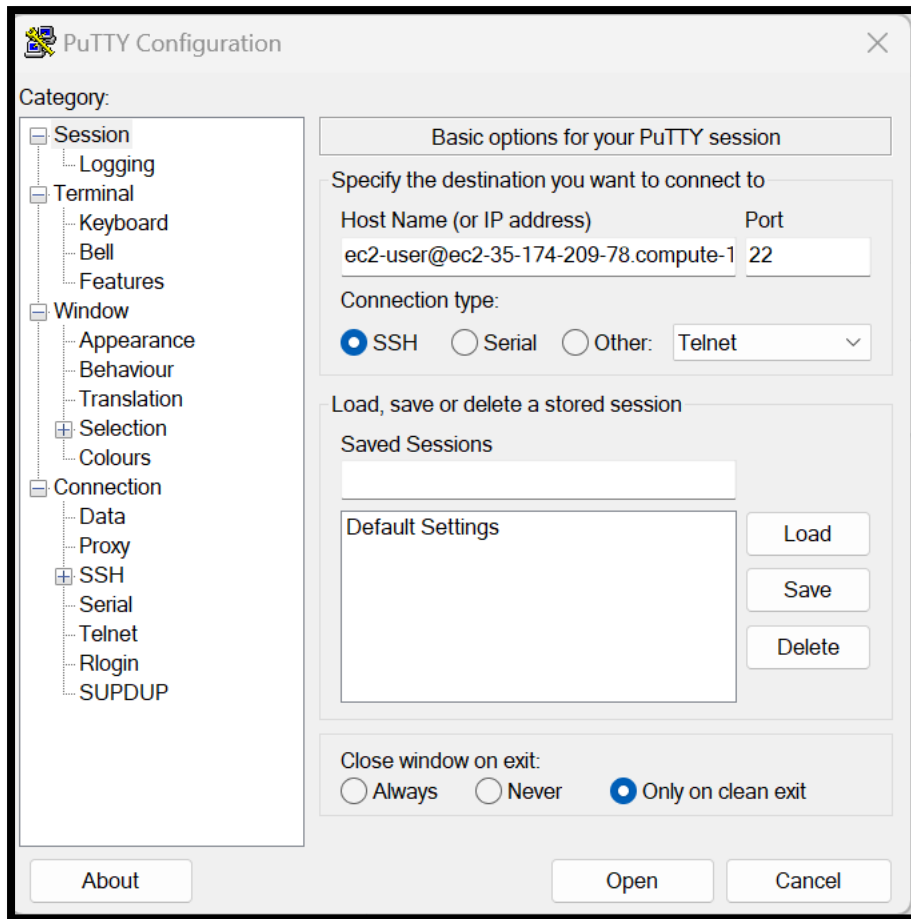
Task 2: Test Auto Scaling

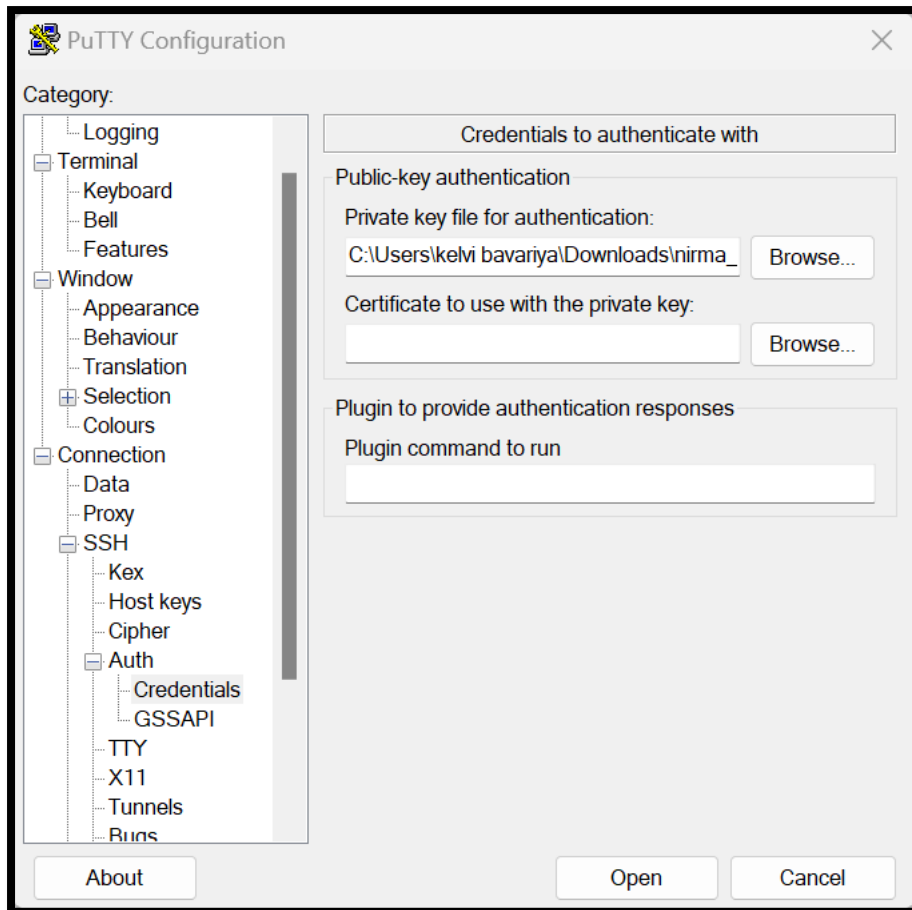
1. Terminating the running instances will create 2 more new instances by autoscaling features.



Instances (5) Info									
Find instance by attribute or tag (case-sensitive)									
<input type="checkbox"/>	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	
<input type="checkbox"/>	-	i-0480b70d5bccb02f0	Terminated	t2.micro	-	No alarms	us-east-1e	-	
<input type="checkbox"/>	-	i-02fe0e10079496527	Running	t2.micro	Initializing	No alarms	us-east-1b	ec2-34-227-176-	
<input type="checkbox"/>	Web Server 1	i-0155e6cde4c2fceb6	Running	t2.micro	2/2 checks passed	No alarms	us-east-1a	-	
<input type="checkbox"/>	Bastion Host	i-05399e678cca95fca	Running	t2.micro	2/2 checks passed	No alarms	us-east-1a	-	
<input type="checkbox"/>	-	i-068250c79710a1c2a	Terminated	t2.micro	-	No alarms	us-east-1c	-	

- Also we can check whether the newly created instances properly work or not .
- For that we will remotely access the website using those newly created instances.





4. Performing same steps as in aws (same commands).

```

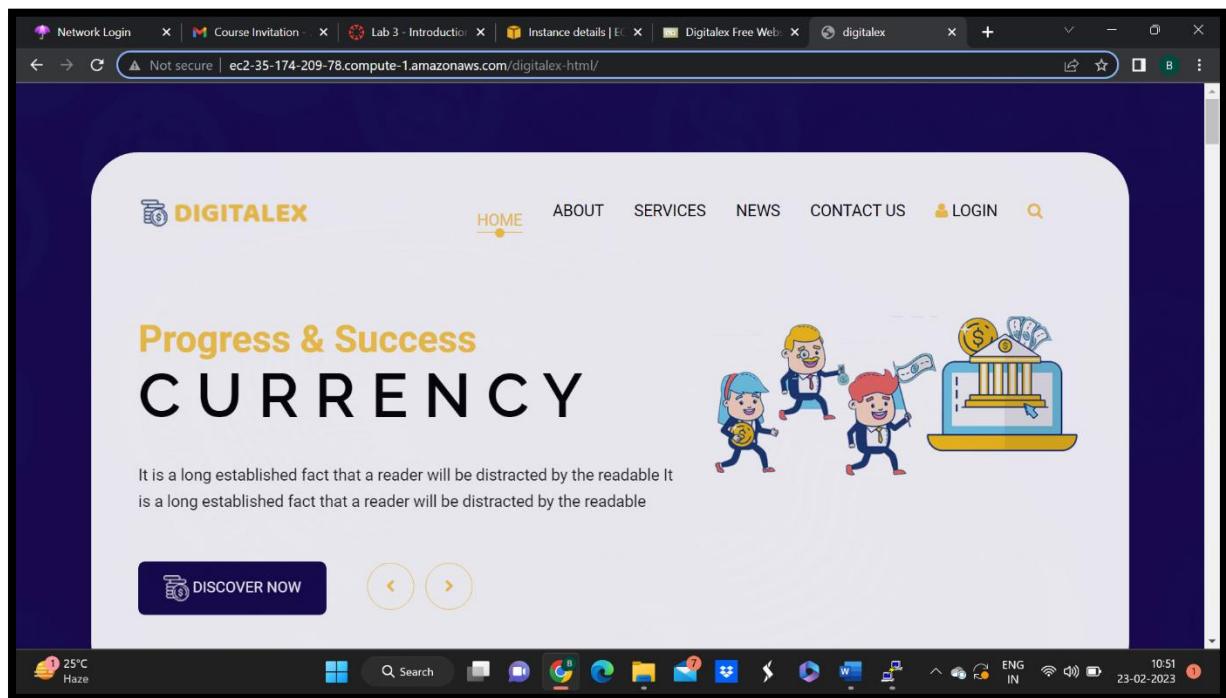
root@ip-172-31-50-65:/var/www/html
Using username "ec2-user".
Authenticating with public key "nirma_key"
Last login: Thu Feb 23 04:52:18 2023 from ec2-18-206-107-27.compute-1.amazonaws.com

 _ _ | _ _ | _ )
 _ | ( _ _ /   Amazon Linux 2 AMI
 _ | \ _ _ | _ |

https://aws.amazon.com/amazon-linux-2/
No packages needed for security; 5 packages available
Run "sudo yum update" to apply all updates.
[ec2-user@ip-172-31-50-65 ~]$ sudo su
[root@ip-172-31-50-65 ec2-user]# yum install update -y
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
amzn2-core                                     | 3.7 kB      00:00
No package update available.
Error: Nothing to do
[root@ip-172-31-50-65 ec2-user]#
[root@ip-172-31-50-65 ec2-user]# yum install httpd
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
Package httpd-2.4.55-1.amzn2.x86_64 already installed and latest version
Nothing to do
[root@ip-172-31-50-65 ec2-user]# systemctl start httpd

```

5. Finally website is hosted using putty software on remote dekstop.



CONCLUSION :-

- AWS Auto Scaling continually monitors your applications to make sure that they are operating at your desired performance levels.
- .When demand spikes, AWS Auto Scaling automatically increases the capacity of constrained resources so you maintain a high quality of service.
- We have learnt how to auto scale the instances for balancing the traffic.
- Also we have verife it by remotely accessing the website.